Project 3 IT3708:

Evolving Neural Networks for a Flatland Agent

Purpose: Learn to implement an evolving neural network to use as a controller for an agent that moves about in a simple, virtual world.

Groups allowed? For this project you may work alone or in groups of two.

1 Assignment

You will use your evolutionary algorithm (EA) to evolve the weights for a small artificial neural network (ANN) whose fitness will depend upon the performance of an agent controlled by the ANN. This agent will do a minimally cognitive task similar to those used in the field of artificial life (ALife) to illustrate the emergence of rudimentary intelligence. Specifically, the agent will move around in a 2-dimensional grid and try to eat as much "food" as possible while avoiding "poison".

2 Flatland: The Testing Environment

The ANNs that you evolve will control an agent that moves about in Flatland (Figure 1), a simple N x N grid environment in which cells contain food, poison or nothing at all. No cell can contain more than one item. The grid is toroidal, meaning that it has no edges nor corners and thus looks like a torus or doughnut. So when the agent moves off the grid to the far right (left), it comes back onto the grid on the far left (right). Similarly, when it moves off the top (bottom) of the grid, it comes back on the bottom (top).

On each timestep, the agent can stay put or move one cell forward, to the its left, or to its right. The agent has a front end from which the directions forward, left and right are based. Before moving left or right, the agent is assumed to turn to the left or right, thus changing its heading direction. So on a left (right) move, the agent actually does 2 things: 1) turns left (right) and 2) moves forward.

When the agent moves into a cell, it automatically eats whatever the cell contains. The agent's goal is to eat as much food as possible, while avoiding the poison.

The Food-Poison Distribution (FPD) for Flatland is a pair of probabilities (f, p), where f is the probability of filling an empty cell with food, while p is the probability of filling one of the remaining unfilled cells with poison. Then, a scenario in Flatland is a grid with food and poison placed in it stochastically, based on a user-chosen FPD. Hence, one FPD can be the basis for many different scenarios.
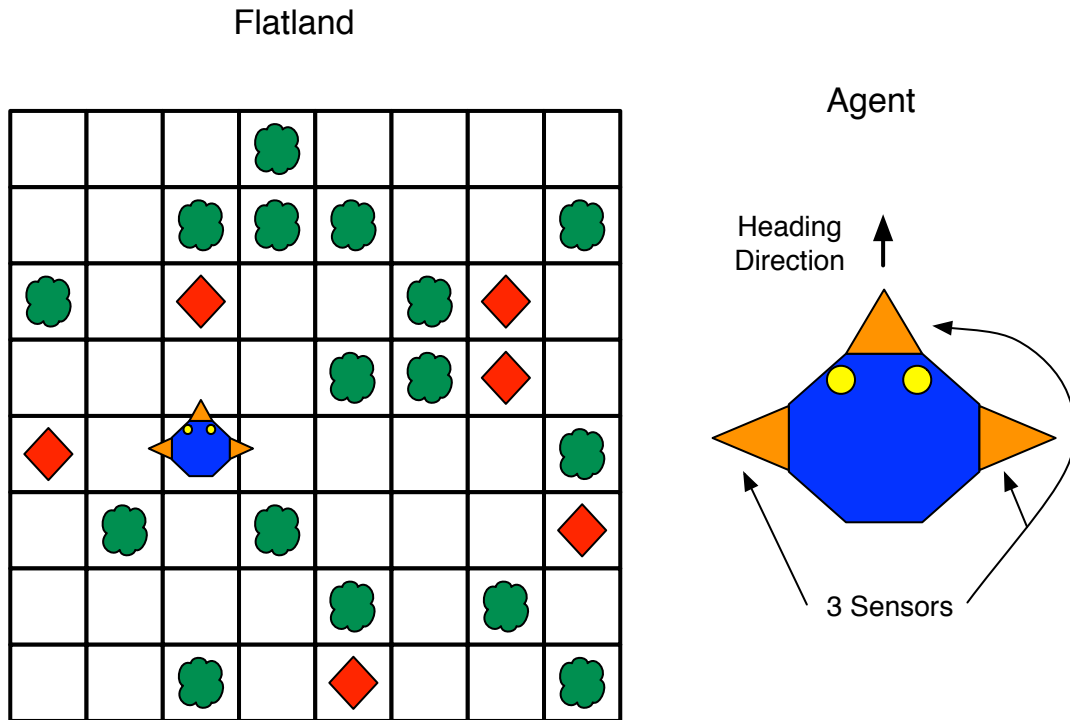
Flatland

Agent

Heading Direction

3 Sensors

Figure 1: Flatland is an N x N grid, with each cell being open or containing one item of food (green blob) or poison (red diamond). The agent has sensors to detect food and poison in the cell immediately in front, to the left, and to the right of itself. No cell contains more than one item, and the agent must eat any item contained within a cell that it visits. Flatland is a toroidal grid: it wraps around from top to bottom and left to right.

For the purposes of this project:

- Flatland is a 10 x 10 grid.

- During fitness testing, each agent will be presented with 1 or 5 scenarios, all based on the same FPD.

- In each generation, all agents are exposed to THE SAME scenario or group of scenarios. However, the scenarios may (or may not) change between generations (as discussed later).

- Each agent gets a total of 60 timesteps per scenario.

- For each scenario, each agent should begin at the same grid cell (of your choice).

- Use an FPD of $(\frac{1}{3}, \frac{1}{3})$ for all scenarios: fill approximately 33% of the cells with food, and then fill approximately 33% *of the remaining cells* with poison.

3 The Artificial Neural Network

This task does not require a complex network such as a CTRNN. Each neuron can simply integrate all weighted inputs and send the sum through an activation function. You should write this ANN yourself.

Figure 2 depicts the 2 input layers and output layer of the ANN. There are separate inputs for sensed food and poison, a design choice that normally makes things easier (though you are free to experiment with a more compact design). The output values of the 3 motor neurons should then be used by your code to determine the agent's movement. You may want to include a firing threshold such that if none of the motor neurons exceed it, then the agent does not move on that timestep.

This is only the skeleton of the ANN; you will need to determine some of the important details, including:

1. The number of (hidden) layers, if any, between the sensory inputs and the motor outputs.

2. The number of neurons in any hidden layers.

3. The activation functions used in the neurons of each layer, including thresholds for functions such as sigmoids, hyperbolic tangents, steps and ramps.

4. Whether or not a bias neuron is needed as input to some or all of the non-input layers.

You can assume that whenever a layer (A) sends connections to a layer (B) that ALL neurons in A get connected to ALL neurons in B.

The weights on all connections must be determined by your evolutionary algorithm (EA). Other parameters can also be evolved, if you so desire, but that is not a requirement.

Be aware that many of the decisions made during ANN design can be highly interdependent. For example, your choice of activation function could affect not only your choice of threshold, but also topological choices, such as the need for and number of hidden nodes. Also, the legal range of weights that you choose can strongly influence the activation functions and thresholds, so you may need to do some simple calculations to determine whether your choices will a) allow neurons to fire, but b) not fire all the time.

The best advice is to begin with a simple ANN and only complexity as needed. Don't start with a 6-layered ANN!

4 The Evolutionary Algorithm

Your EA need not be very complex. The genotype can be a simple bit vector, while the phenotype can be a list of connection weights. Mutation then involves bit flips, while crossover is the traditional version for bit vectors. Some textbooks recommend special crossover operators for ANN weight vectors, but these are not necessary for this assignment.

The evolutionary runs that you perform will be of two types:

- **Dynamic** - the set of randomly-generated scenarios changes after each generation. So each new generation of individuals experiences a new set of scenarios.

Figure 2: The basic structure of the ANN for a Flatland agent: 6 sensory neurons, 3 each for food or poison in the forward, left or right cell; and 3 motor neurons for the 3 possible directions of movement: forward, left or right.

- **Static** - the same randomly-generated scenarios are used throughout the ENTIRE evolutionary run, i.e., in each and every generation.

All of your evolutionary runs should include some degree of **elitism**: the best individual(s) from generation K are directly copied into the population of generation K+1. Decide for yourself whether more than one such individual needs to be copied.

Other parameters such as population size, number of generations, mutation and crossover rates, adult selection strategies and parent selection mechanisms, etc. will also need to be determined experimentally: try combinations until you find one that gives good results. **Note**: what worked well for previous problems may not work well for this problem.

4.1 The Fitness Function

This is another detail that you will have to figure out yourself. One of the key issues is the amount of penalty associated with eating poison. There is a wide spectrum, from killing (i.e. giving a fitness of zero to) any agent that ingests any poison, to more forgiving approaches that simply detract a small, fixed amount from total fitness for each poison item consumed. Too weak a

penalty will not encourage poison-avoidance at all, while too harsh a penalty can produce agents that do nothing (i.e. hardly move at all or walk in circles), for fear of making a fatal mistake.

4.2 Expected Performance

For this task, it can be difficult to evolve an optimal controller, but your EA should be able to find weight vectors that produce reasonably intelligent behavior in an agent: it should be able to eat most of the food and avoid most of the poison in all 5 60-step scenarios (for both static and dynamic runs). Do not expect perfection, but if you are NOT getting a very clear sense that your agent can differentiate between food and poison in MOST contexts (i.e. combinations of food and poison inputs to the 3 sensors), then you and your EA have more work to do.

4.3 Visualization

At the end of an evolutionary run, you must be able to visualize the behavior of the best-of-run individual on the scenarios. For static runs, these will be the same scenarios that the agent was originally tested on, but it should be possible to also give the agent new, randomly-generated scenarios as well. For dynamic runs, these will always be new, randomly-generated scenarios.

This must be a complete 2-d visualization of a grid, not a command-line printout of an array. You must use different colors for food, poison and the agent so that it is very easy to immediately see the structure of a scenario and the behavior of the agent. It should be easy to see the direction of the agent as well. The delay between visualized timesteps must be adjustable so that we can slow down the run to see exactly what the agent is doing.

You should also be able to visualize the EA run on a plot, as outlined in the previous project description. *For this project, please show one run in each plot, not averages over multiple runs.*

5 Report

You should write a report answering the points below. The report can give a maximum score of 12 points. Your report must not exceed 4 pages in total. Overlength reports will result in points being deducted from your final score. Print on both sides of the sheets, preferably. Bring a hard copy of your report to the demo session

a) Document your EA choices. (3p)

- A table summarizing the main EA parameters used for the successful runs of your system, along with a brief description of the exploratory process that you used to find them..

- A complete description of your fitness function in terms of both a mathematical expression and explanatory text.

b) Document your implementation (3p)

- A detailed description of the complete ANN design (excluding the actual connection-weight values) that proved most successful in your experiments. This includes the layers

of neurons, their activation functions, any thresholds or biases, and the range of acceptable weights that your EA had to work with. You do not need to describe all of the individual weights.

- Briefly describe the process that you went through to find your design. Describe why your design should be able to solve the problem.

c) Performance of the EA (6p)

- Do a **static** run where each agent is only tested on a single scenario each generation. Include a plot of the result. Describe the behavior of the best evolved agent on this scenario. Also test the evolved agent on a new, randomly-generated scenario and explain the behavior of it there.

- Do the same for a new **static** run, but this time have 5 different scenarios to test each agent with each generation.

- Do a **dynamic** run where each agent is only tested on a single scenario. Include a plot in the result. Describe the behavior of the best evolved agent on a new, randomly-generated scenario.

- Do the same for a new **dynamic** run, this time with 5 different scenarios to test each agent.

- Do you see any difference between the behavior of the agents in the four cases above? Discuss why (not). Also discuss any differences between the plots.

6 Demo

There will be a demo session where you will show us the running code and we will verify that it works. This demonstration can give a maximum of 8 points.

For this demo, you will use your EA to evolve agents and show us that they behave reasonably intelligent. Your EA should be able to produce these good agents in just a few minutes during the demo.

7 Delivery

You should deliver your report + a zip file of your code on It's Learning. The deadline is given on the assignment on It's Learning. The 20 points total for this project is 20 of the 100 points available for this class. For this project you can work alone or in groups of two.

If you work in a group, you only need to deliver once on It's Learning (but both group members must the registered as part of the submission on It's Learning!). Both group members need to attend the demo session.